

# Messaging Engine Documentation

## PMM Messaging Engine — Technical Documentation

*Comprehensive Architecture, Pipeline, and Prompt Reference*

### 1. Executive Summary

The PMM Messaging Engine converts product documentation into scored, quality-tested messaging assets through 5 specialized AI pipelines. It follows an “outside-in” philosophy — starting with real practitioner pain from developer communities rather than vendor features.

### Key Capabilities

- **5 distinct generation pipelines** — Standard, Outside-In (signature), Adversarial, Multi-Perspective, and Straight Through
- **Sequential DAG architecture** — each pipeline step feeds the next
- **Shared refinement loop** — score → deslop → refine (up to 3x) with plateau detection
- **Multi-model AI strategy** — Gemini 3 Pro/Flash (production) with test profile fallback
- **5-dimension quality scoring** with scorer health tracking and degraded mode
- **8 asset types** from battlecards to narratives
- **Voice profile system** with per-voice quality thresholds

### Tech Stack

Layer	Technology
Backend	Hono (REST API), TypeScript
Frontend	Vite + React + Tailwind CSS
Database	SQLite + Drizzle ORM (14 tables)
AI Models	Gemini 3 Pro/Flash (prod), Gemini 2.5 Flash (test), Claude (opt-in)

### 2. System Architecture

#### File Structure

```
src/
  └── api/
    ├── generate.ts          # ~186 lines – route definitions +
    │   re-exports only
    │   ├── validation.ts     # Zod schemas for all endpoints
    │   │   ├── index.ts       # Hono app + route mounting
    │   │   └── middleware/
    │   │       ├── auth.ts    # JWT authentication (hard error)
    │   │       └── rate-limit.ts # Per-IP HTTP rate limiting
    │   └── admin/
    │       ├── documents      # Admin CRUD (settings, voices,
    │       └── workspace/       # Session management, chat, actions
    ├── config.ts              # Model profiles, env config
    └── types/index.ts         # Shared interfaces
        (ScoringThresholds, PipelineStep, etc.)
```

```

├── services/
│   ├── pipeline/
│   │   ├── orchestrator.ts      # Job lifecycle, dispatch,
│   │   └── generateAndScore, refinementLoop, storeVariant
│   ├── prompts.ts              # All prompt builders, templates,
│   └── banned-words
│       └── evidence.ts        # EvidenceBundle,
community/competitive-research
    └── pipelines/
        ├── standard.ts        # PoV-first with deep extraction
        ├── outside-in.ts       # Community pain-first (signature)
        ├── adversarial.ts      # Attack/defend generation
        ├── multi-perspective.ts # 3 angles + synthesis
        └── straight-through.ts  # Score-only, no generation
    └── quality/
        └── score-content.ts    # 5 parallel scorers + health
tracking
    ├── slop-detector.ts       # AI cliché detection
    ├── vendor-speak.ts        # Marketing jargon scorer
    ├── authenticity.ts        # Authenticity scorer
    ├── specificity.ts         # Specificity scorer
    ├── persona-critic.ts      # Persona fit scorer
    └── grounding-validator.ts # Strips fabricated quotes
    ├── ai/clients.ts          # Gemini + Claude API wrappers
    ├── product/insights.ts    # Insight extraction + deep PoV
    └── research/deep-research.ts # Gemini Deep Research agent
db/
utils/                      # Schema, connection, seed data
                            # Logger, hash, retry

```

## Security

Layer	Mechanism
Authentication	JWT — hard error on invalid/missing token
Validation	Zod schemas on all request bodies
Rate Limiting	Per-IP, per-path (Flash: 60/min, Pro: 15/min)
Input Limits	productDocs: 500K chars, prompt: 10K chars

## 3. Model Profile System

Controlled by MODEL\_PROFILE env var. Every model call logs the actual model name and emits it via pipeline step events for UI visibility.

Task	Production	Test
flash	gemini-3-flash-preview	gemini-2.5-flash
pro	gemini-3-pro-preview	gemini-2.5-flash
deepResearch	deep-research-pro-preview-12-2025	gemini-2.5-flash
generation	gemini-3-pro-preview	gemini-2.5-flash
scoring	gemini-3-flash-preview	gemini-2.5-flash
deslop	gemini-3-pro-preview	gemini-2.5-flash

Test profile uses a single cheap model for all tasks — ideal for development and CI.

## 4. Quality Scoring System

5 independent scorers run in parallel via Promise.all:

Scorer	Scale	Direction	What It Measures
Slop	0-10	Lower is better	AI clichés, filler phrases

Vendor Speak	0-10	Lower is better	Marketing jargon, buzzwords
Authenticity	0-10	Higher is better	Genuine practitioner voice
Specificity	0-10	Higher is better	Concrete details vs vague claims
Persona	0-10	Higher is better	Fit with target persona

## Scorer Health & Degraded Mode

Each scoring run tracks health: { succeeded: N, failed: [...], total: 5 }.

If a scorer throws, it returns a neutral fallback score of 5 and the failure is logged. The system continues in degraded mode rather than failing the entire job. Health is visible in the pipeline step UI.

## Quality Gates (Per-Voice Thresholds)

Gate	Default	Direction
slopMax	5	Reject if above
vendorSpeakMax	5	Reject if above
authenticityMin	6	Reject if below
specificityMin	6	Reject if below
personaMin	6	Reject if below

## 5. Pipelines

All pipelines follow a sequential DAG pattern. Each step feeds the next — no branching.

### 5.1 Standard Pipeline

**PoV-first — extract a deep product thesis before generating.**

```
Deep PoV Extraction (Pro)
  → Insight Extraction (Flash)
  → Community Research (Deep Research)
  → Competitive Research (Deep Research)
  → For each assetType × voice:
    → Generate with PoV-first prompt (Pro)
    → Score → Refinement Loop (up to 3x) → Store
```

### 5.2 Outside-In Pipeline (Signature)

**Community pain-first — start with real practitioner frustrations.**

```
Insight Extraction (Flash)
  → Community Deep Research (Deep Research)
  → Competitive Research (Deep Research)
  → For each assetType × voice:
    → Generate with pain-first prompt (Pro)
    → Score → Refinement Loop (up to 3x) → Store
```

### 5.3 Adversarial Pipeline

**Attack/defend — generate objections, then craft messaging that survives them.**

```
Insight Extraction (Flash)
  → Research (Deep Research)
  → For each assetType × voice:
    → Generate attack (hostile buyer objections)
    → Generate defense (messaging addressing attacks)
```

## 5.4 Multi-Perspective Pipeline

**3 angles + synthesis — practitioner, buyer, executive perspectives merged.**

```
Insight Extraction (Flash)
  → Research (Deep Research)
  → For each assetType × voice:
    → Generate Practitioner angle
    → Generate Buyer angle
    → Generate Executive angle
    → Synthesize into unified content
  → Score → Refinement Loop (up to 3x) → Store
```

## 5.5 Straight Through Pipeline

**Score-only — evaluate existing content without any generation.**

```
Insight Extraction (Flash)
  → For each assetType × voice:
    → Score existing content (5 scorers)
    → Store (original content preserved as-is)
```

**Key differences:** - NO generation — content is never rewritten - NO refinement loop — no deslop, no iteration - NO research phase - Requires existingMessaging — fails if not provided

**Use case:** Benchmarking existing marketing content.

---

## 6. Shared Pipeline Components

### Refinement Loop

1. Score with 5 scorers
2. If quality gates pass → done
3. Deslop to remove AI clichés
4. Build refinement prompt with score feedback
5. Regenerate
6. Rescore — stop if plateau detected
7. Repeat up to 3 times

### Evidence Bundle

Community and competitive research packaged as EvidenceBundle: - communityPain — real developer frustrations - competitorWeaknesses — positioning gaps - practitionerQuotes — attributed quotes with source URLs - evidenceLevel — strong | moderate | weak

---

## 7. API Reference

### Core Endpoints

Method	Path	Auth	Description
POST	/api/generate	JWT	Start generation job
GET	/api/voices	JWT	List voice profiles
POST	/api/extract	JWT	Extract document content
POST	/api/auth/login	—	Get JWT token
POST	/api/auth/signup	—	Create account

### Generate Request Schema

```
{
  "productDocs": "string (required, max 500K)",
```

```

    "existingMessaging": "string (optional, max 500K)",
    "prompt": "string (optional, max 10K)",
    "voiceProfileIds": ["string (min 1)"],
    "assetTypes": ["battlecard | talk_track | launch.messaging | social_hook | one_pager | email_copy | messaging_template | narrative"],
    "model": "string (optional)",
    "pipeline": "standard | outside-in | adversarial | multi-perspective | straight-through"
}

```

## Asset Types

Type	Label
battlecard	Battlecard
talk_track	Talk Track
launch.messaging	Launch Messaging
social_hook	Social Hook
one_pager	One-Pager
email_copy	Email Copy
messaging_template	Messaging Template
narrative	Narrative