

PMM Messaging Engine

PMM Messaging Engine — Technical Documentation

Comprehensive Architecture, Pipeline, and Prompt Reference

1. Executive Summary

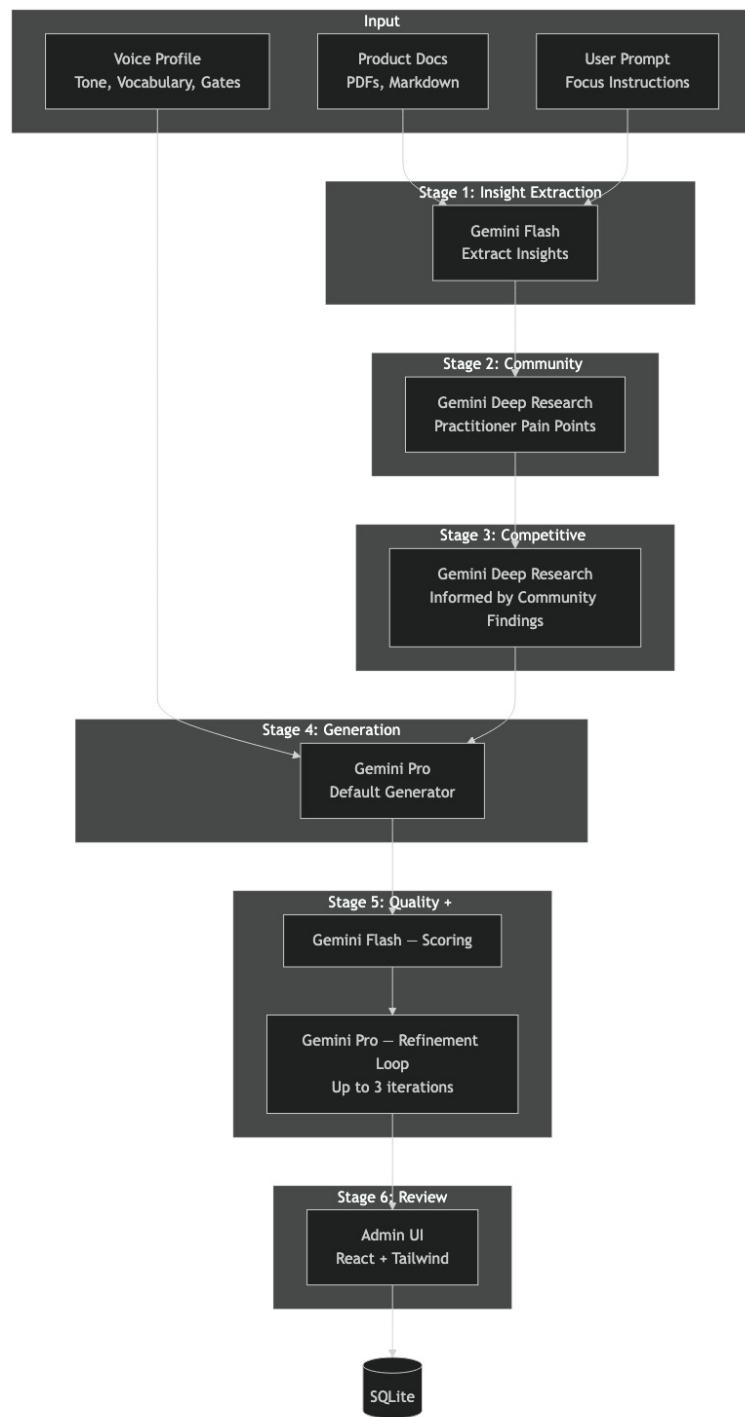
The PMM Messaging Engine converts product documentation into scored, quality-tested messaging assets through 5 specialized AI pipelines. It follows an “outside-in” philosophy — starting with real practitioner pain from developer communities rather than vendor features.

Key Capabilities

- **5 distinct generation pipelines** — Standard, Outside-In (signature), Adversarial, Multi-Perspective, and Straight Through
- **Sequential DAG architecture** — each pipeline step feeds the next
- **Shared refinement loop** — score → deslop → refine (up to 3x) with plateau detection
- **Multi-model AI strategy** — Gemini 3 Pro/Flash (production) with test profile fallback
- **5-dimension quality scoring** with scorer health tracking and degraded mode
- **8 asset types** from battlecards to narratives
- **Voice profile system** with per-voice quality thresholds

Tech Stack

Layer	Technology
Backend	Hono (REST API), TypeScript
Frontend	Vite + React + Tailwind CSS
Database	SQLite + Drizzle ORM (14 tables)
AI Models	Gemini 3 Pro/Flash (prod), Gemini 2.5 Flash (test), Claude (opt-in)



System Architecture

2. System Architecture

File Structure

```

src/
  └── api/
      ├── generate.ts          # ~186 lines – route definitions +
      |   re-exports only
      |   ├── validation.ts     # Zod schemas for all endpoints
      |   ├── index.ts          # Hono app + route mounting
      |   └── middleware/
      |       ├── auth.ts        # JWT authentication (hard error)
      |       └── rate-limit.ts  # Per-IP HTTP rate limiting
      └── admin/
          documents)           # Admin CRUD (settings, voices,
          └── workspace/         # Session management, chat, actions
          config.ts             # Model profiles, env config
          types/index.ts        # Shared interfaces
          (ScoringThresholds, PipelineStep, etc.)
  └── services/
      └── pipeline/

```

```

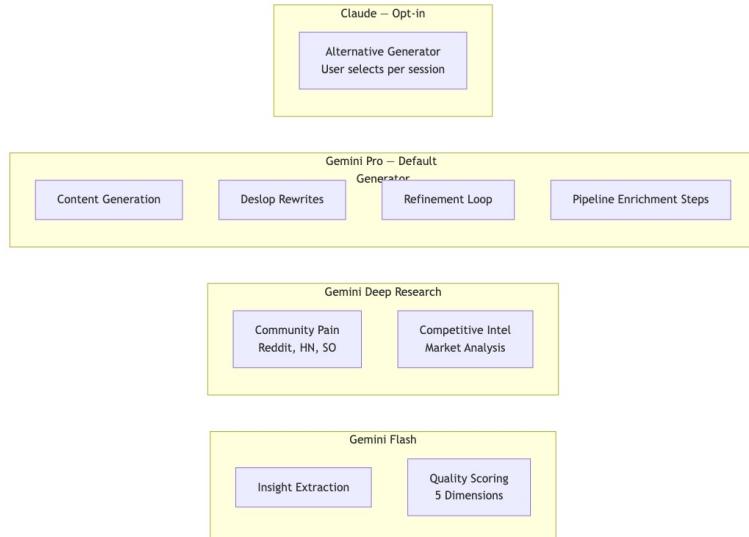
|   |   └── orchestrator.ts      # Job lifecycle, dispatch,
generateAndScore, refinementLoop, storeVariant
|   |   └── prompts.ts          # All prompt builders, templates,
banned words
|   |       └── evidence.ts      # EvidenceBundle,
community/competitive research
|   |           └── pipelines/
|   |               ├── standard.ts    # PoV-first with deep extraction
|   |               ├── outside-in.ts  # Community pain-first (signature)
|   |               ├── adversarial.ts # Attack/defend generation
|   |               ├── multi-perspective.ts # 3 angles + synthesis
|   |               └── straight-through.ts # Score-only, no generation
|   └── quality/
|       └── score-content.ts    # 5 parallel scorers + health
tracking
|           └── slop-detector.ts    # AI cliché detection
|           └── vendor-speak.ts     # Marketing jargon scorer
|           └── authenticity.ts    # Authenticity scorer
|           └── specificity.ts     # Specificity scorer
|           └── persona-critic.ts   # Persona fit scorer
|               └── grounding-validator.ts # Strips fabricated quotes
|           └── ai/clients.ts        # Gemini + Claude API wrappers
|           └── product/insights.ts   # Insight extraction + deep PoV
|           └── research/deep-research.ts # Gemini Deep Research agent
└── db/                                # Schema, connection, seed data
└── utils/                             # Logger, hash, retry

```

Security

Layer	Mechanism
Authentication	JWT — hard error on invalid/missing token
Validation	Zod schemas on all request bodies
Rate Limiting	Per-IP, per-path (Flash: 60/min, Pro: 15/min)
Input Limits	productDocs: 500K chars, prompt: 10K chars

3. Model Profile System



Multi-Model Strategy

Controlled by `MODEL_PROFILE` env var. Every model call logs the actual model name and emits it via pipeline step events for UI visibility.

Task	Production	Test
flash	gemini-3-flash-preview	gemini-2.5-flash
pro	gemini-3-pro-preview	gemini-2.5-flash
deepResearch	deep-research-pro-preview-12-2025	gemini-2.5-flash

generation	gemini-3-pro-preview	gemini-2.5-flash
scoring	gemini-3-flash-preview	gemini-2.5-flash
deslop	gemini-3-pro-preview	gemini-2.5-flash

Test profile uses a single cheap model for all tasks — ideal for development and CI.

4. Quality Scoring System

5 independent scorers run in parallel via `Promise.all`:

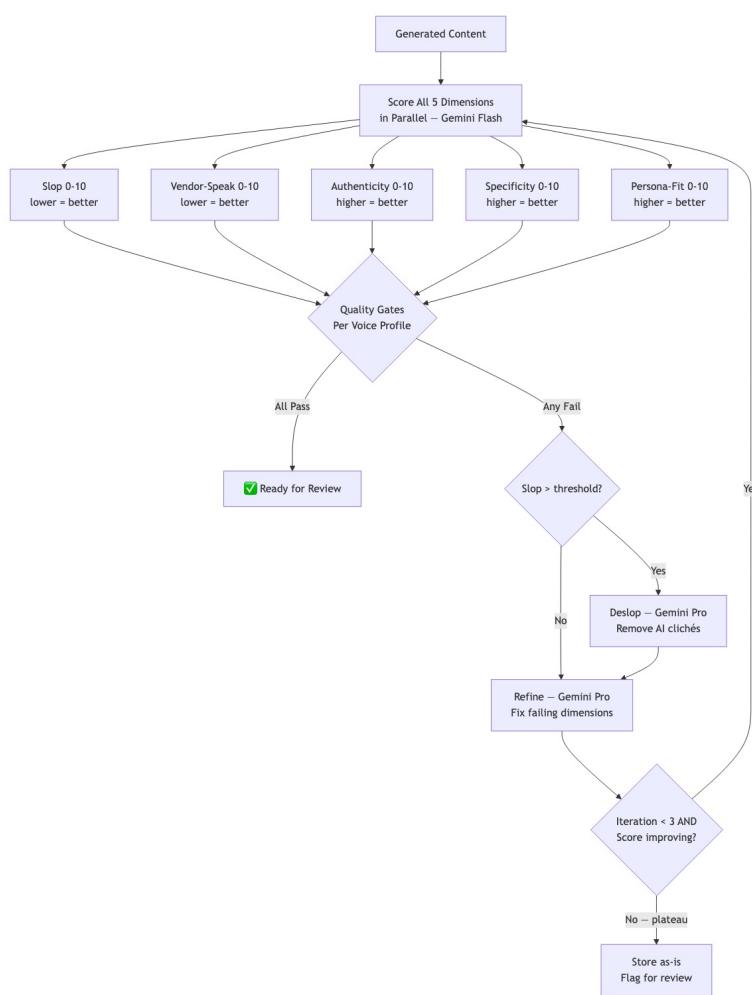
Scorer	Scale	Direction	What It Measures
Slop	0-10	Lower is better	AI clichés, filler phrases
Vendor Speak	0-10	Lower is better	Marketing jargon, buzzwords
Authenticity	0-10	Higher is better	Genuine practitioner voice
Specificity	0-10	Higher is better	Concrete details vs vague claims
Persona	0-10	Higher is better	Fit with target persona

Scorer Health & Degraded Mode

Each scoring run tracks health: { succeeded: N, failed: [...], total: 5 }.

If a scorer throws, it returns a neutral fallback score of 5 and the failure is logged. The system continues in degraded mode rather than failing the entire job. Health is visible in the pipeline step UI.

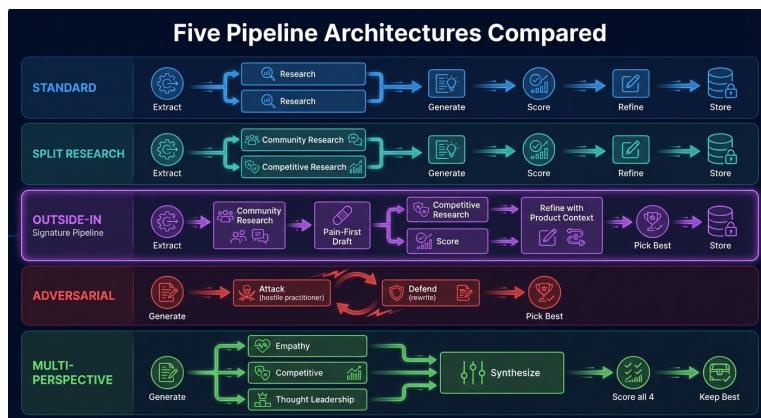
Quality Gates (Per-Voice Thresholds)



Quality Pipeline

Gate	Default	Direction
slopMax	5	Reject if above
vendorSpeakMax	5	Reject if above
authenticityMin	6	Reject if below
specificityMin	6	Reject if below
personaMin	6	Reject if below

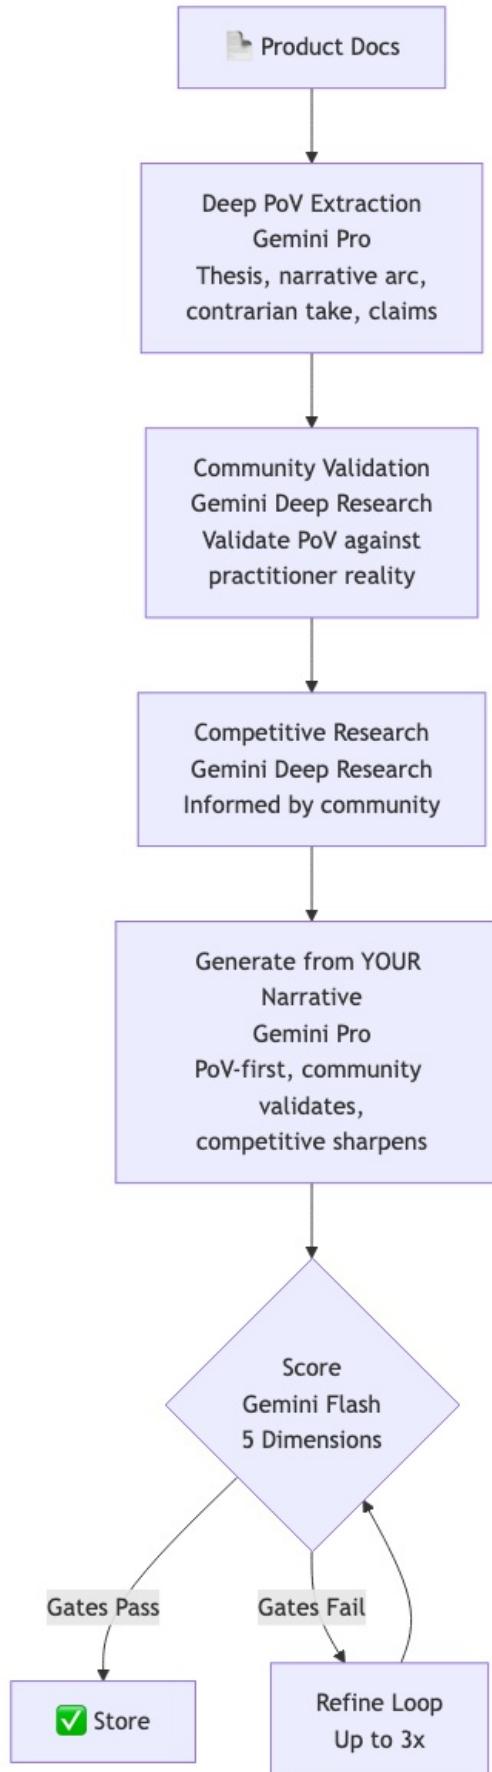
5. Pipelines



All Pipelines

All pipelines follow a sequential DAG pattern. Each step feeds the next — no branching.

5.1 Standard Pipeline

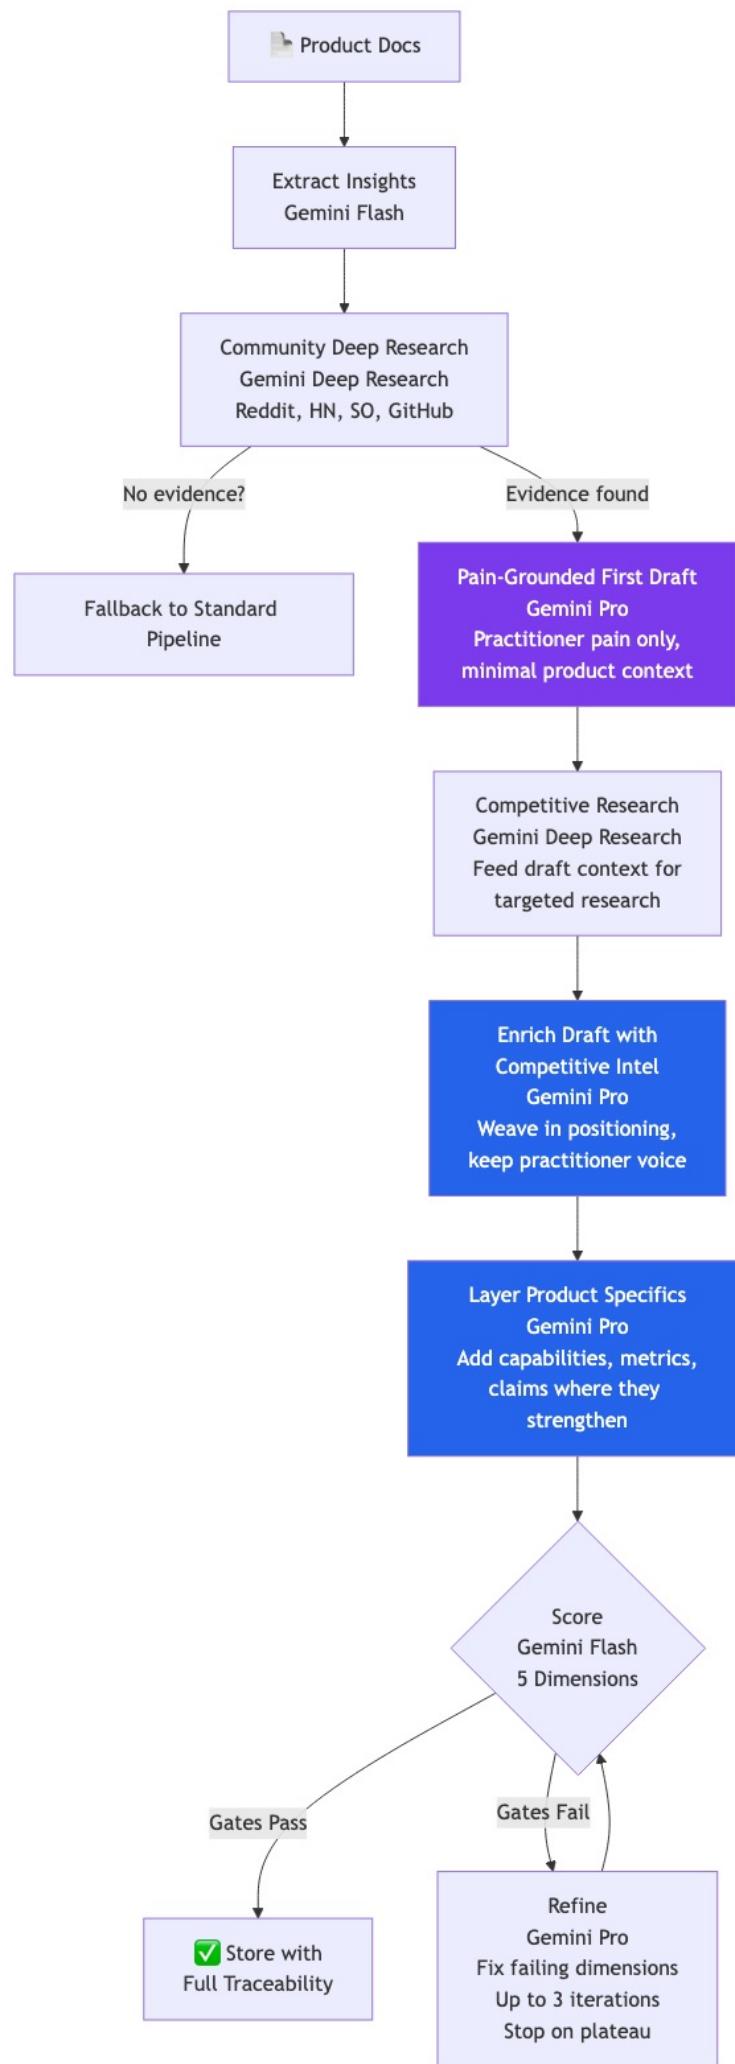


Standard Pipeline

PoV-first — extract a deep product thesis before generating.

- Deep PoV Extraction (Pro)
 - Insight Extraction (Flash)
 - Community Research (Deep Research)
 - Competitive Research (Deep Research)
 - For each assetType × voice:
 - Generate with PoV-first prompt (Pro)
 - Score → Refinement Loop (up to 3x) → Store

5.2 Outside-In Pipeline (Signature)



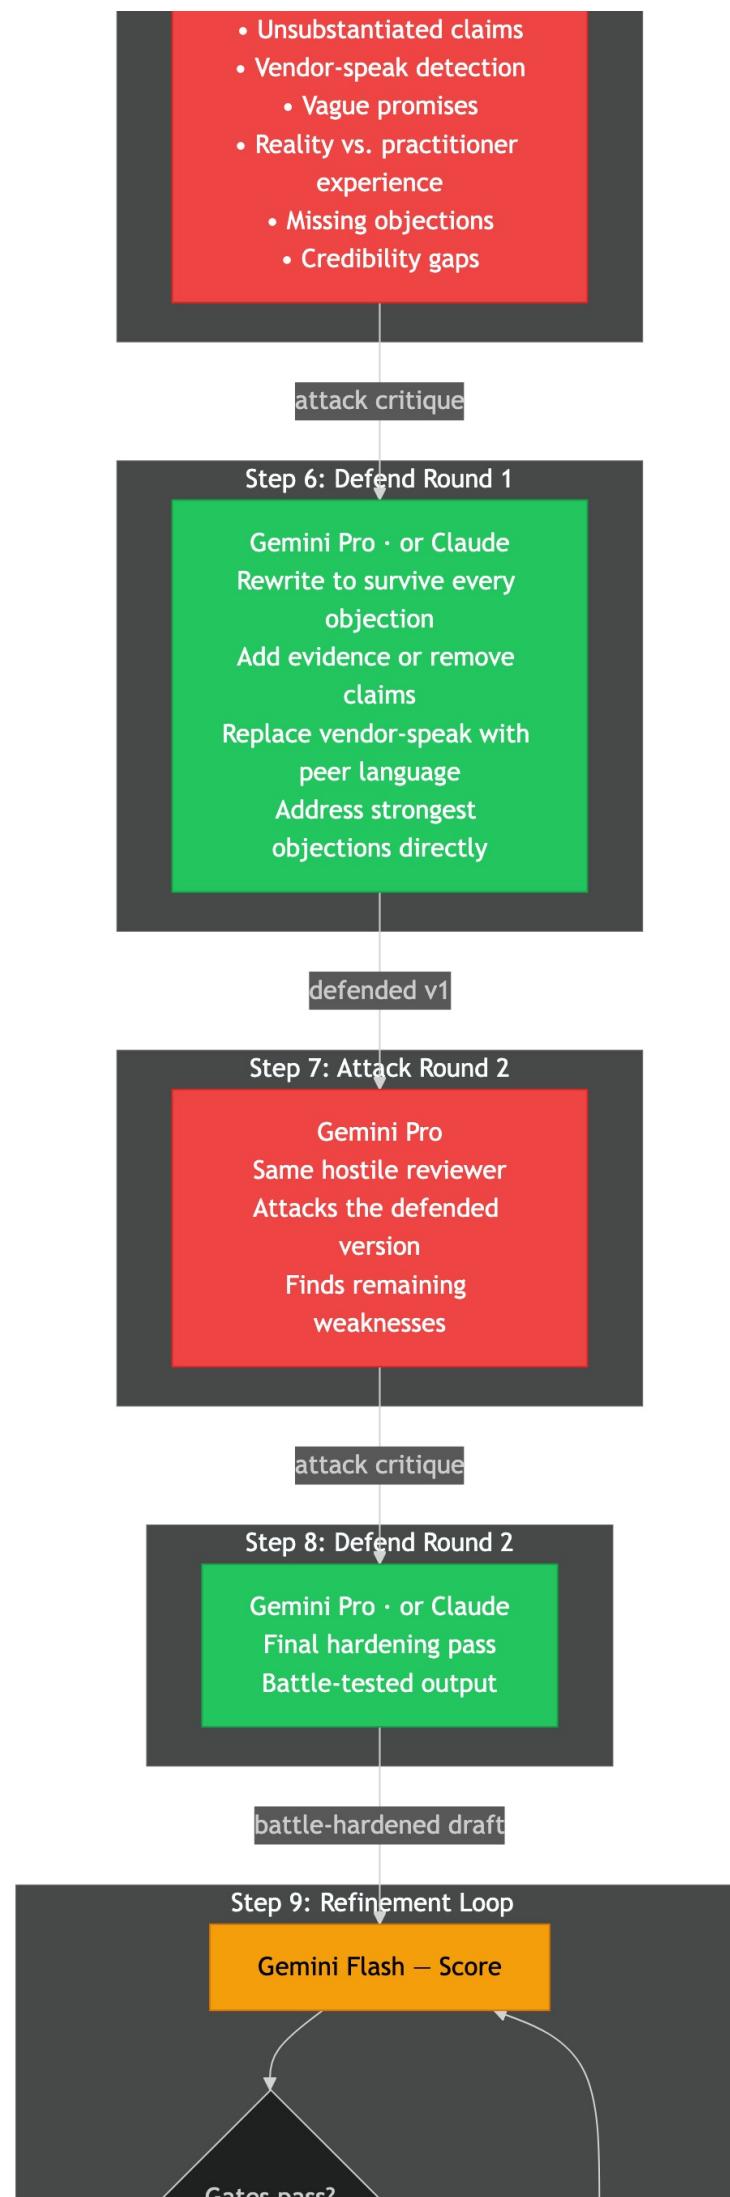
Outside-In Pipeline

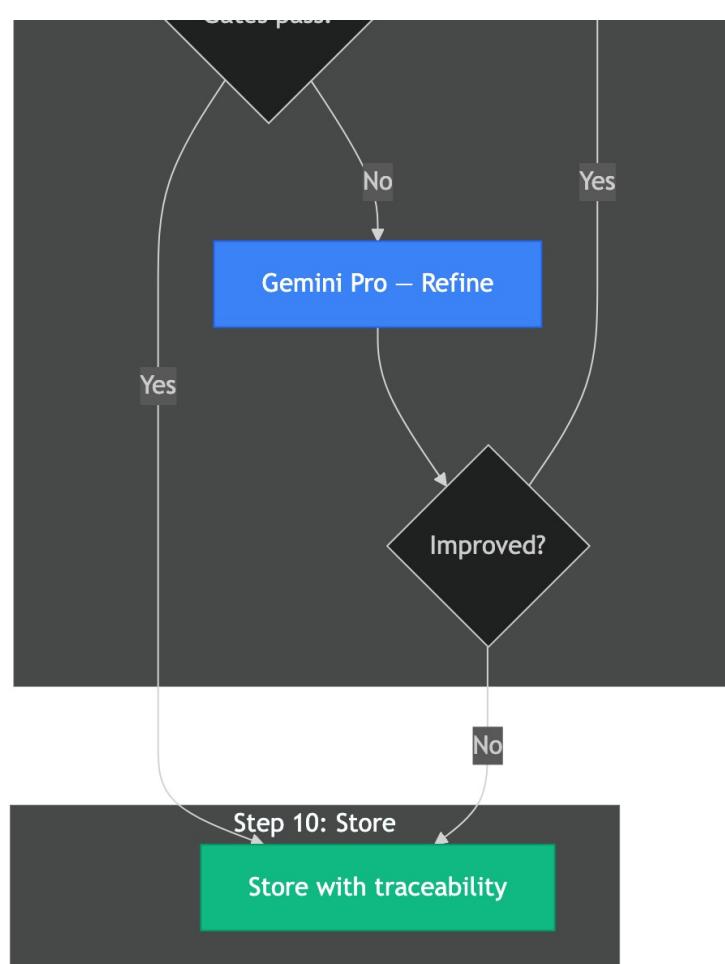
Community pain-first — start with real practitioner frustrations.

Insight Extraction (Flash)
→ Community Deep Research (Deep Research)
→ Competitive Research (Deep Research)
→ For each assetType x voice:
 → Generate with pain-first prompt (Pro)
 → Score → Refinement Loop (up to 3x) → Store

5.3 Adversarial Pipeline







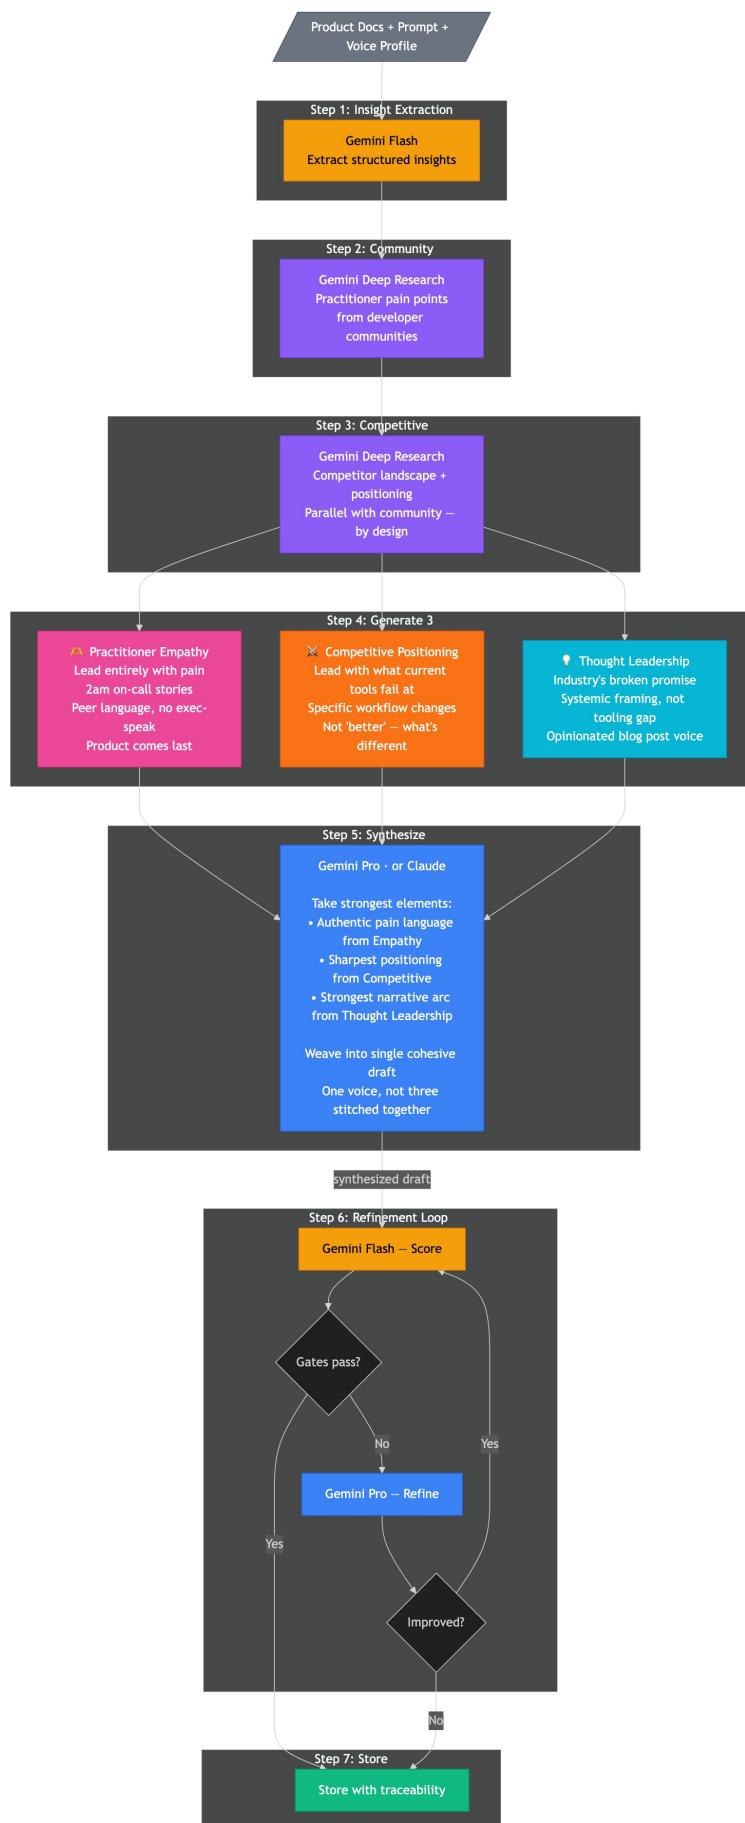
Adversarial Pipeline

Attack/defend — generate objections, then craft messaging that survives them.

Insight Extraction (Flash)

- Research (Deep Research)
- For each assetType × voice:
 - Generate attack (hostile buyer objections)
 - Generate defense (messaging addressing attacks)
 - Score defense → Refinement Loop (up to 3x) → Store

5.4 Multi-Perspective Pipeline



Multi-Perspective Pipeline

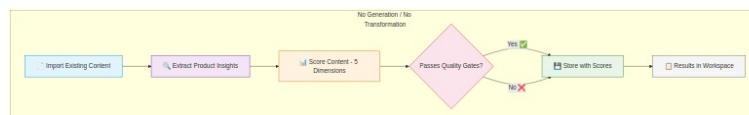
3 angles + synthesis — practitioner, buyer, executive perspectives merged.

Insight Extraction (Flash)

- Research (Deep Research)
- For each assetType x voice:
 - Generate Practitioner angle
 - Generate Buyer angle
 - Generate Executive angle

- Synthesize into unified content
- Score → Refinement Loop (up to 3x) → Store

5.5 Straight Through Pipeline



Straight Through Pipeline

Score-only – evaluate existing content without any generation.

Insight Extraction (Flash)

- For each assetType × voice:
- Score existing content (5 scorers)
- Store (original content preserved as-is)

Key differences: - NO generation — content is never rewritten - NO refinement loop — no deslop, no iteration - NO research phase - Requires existingMessaging — fails if not provided

Use case: Benchmarking existing marketing content.

6. Shared Pipeline Components

Refinement Loop

1. Score with 5 scorers
2. If quality gates pass → done
3. Deslop to remove AI clichés
4. Build refinement prompt with score feedback
5. Regenerate
6. Rescore — stop if plateau detected
7. Repeat up to 3 times

Evidence Bundle

Community and competitive research packaged as EvidenceBundle: - communityPain — real developer frustrations - competitorWeaknesses — positioning gaps - practitionerQuotes — attributed quotes with source URLs - evidenceLevel — strong | moderate | weak

7. API Reference

Core Endpoints

Method	Path	Auth	Description
POST	/api/generate	JWT	Start generation job
GET	/api/voices	JWT	List voice profiles
POST	/api/extract	JWT	Extract document content
POST	/api/auth/login	—	Get JWT token
POST	/api/auth/signup	—	Create account

Generate Request Schema

```
{
  "productDocs": "string (required, max 500K)",
  "existingMessaging": "string (optional, max 500K)",
  "prompt": "string (optional, max 10K)",
  "voiceProfileIds": ["string (min 1)"],
  "assetTypes": ["battlecard | talk_track | launch.messaging | social_hook | one_pager | email_copy | messaging_template | narrative"],
  "model": "string (optional)",
}
```

```
    "pipeline": "standard | outside-in | adversarial | multi-  
perspective | straight-through"  
}
```

Asset Types

Type	Label
battlecard	Battlecard
talk_track	Talk Track
launch.messaging	Launch Messaging
social_hook	Social Hook
one_pager	One-Pager
email_copy	Email Copy
messaging_template	Messaging Template
narrative	Narrative